

WHEN SIMPLE API CALL IS
NOT ENOUGH

OUTBOX PATTERN

JACEK MILEWSKI

✉ jacek.milewski.k@gmail.com

🐦 @jacek_mil



AGENDA

1. THE PROBLEM
2. THE OUTBOX PATTERN
3. IMPLEMENTATION
4. TESTING
5. CONCERNS


THE PROBLEM

THE PROBLEM




DETAILS


SAVE IT IN DB
AND
POST A REQUEST TO EXTERNAL API



SYNC? ASYNC?

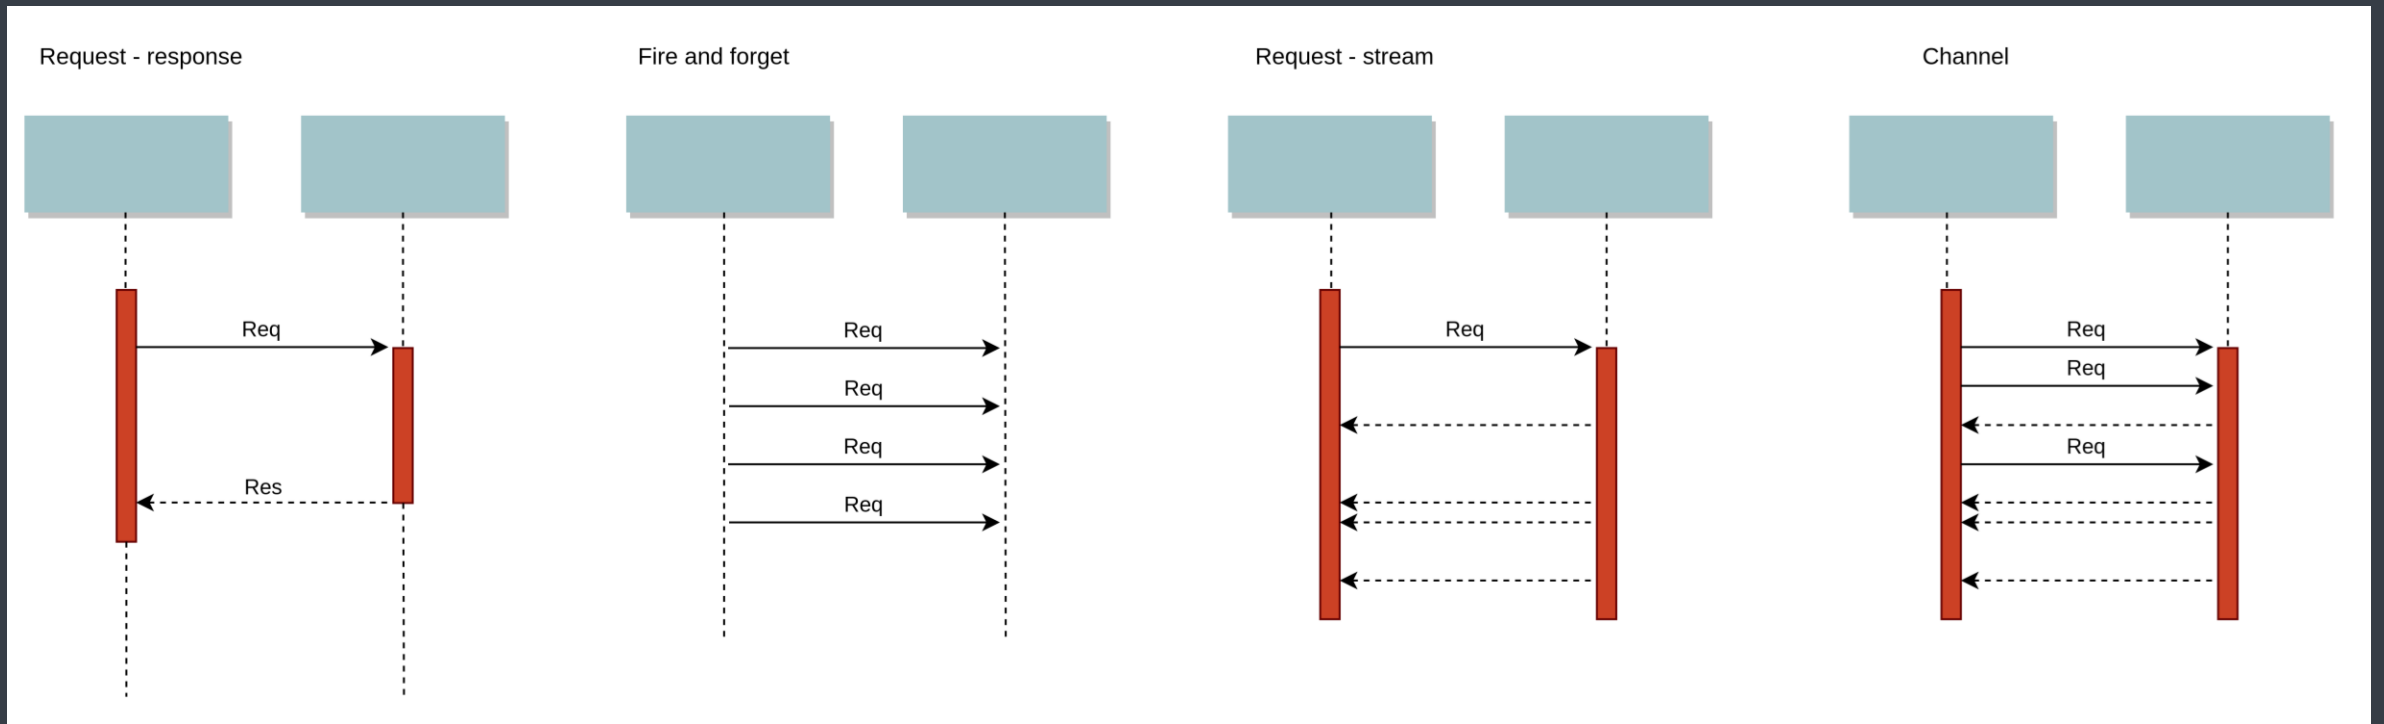


SYNC? ASYNC?



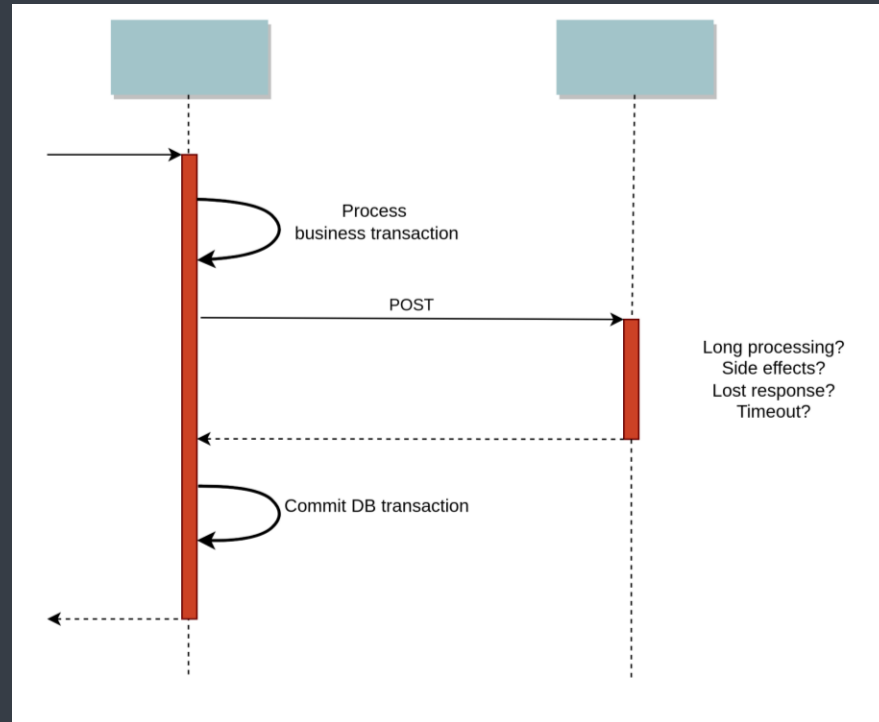
	Sync	async
consistency	Immediate	eventual
Implementation effort	Small	Complex
Processing time	Slower	Same but split to steps
Handling errors / rollbacks	Explicit error responses	Reconciliation needed
Retry mechanism	None	Simpler, depending on a tools

STYLES OF COMMUNICATION

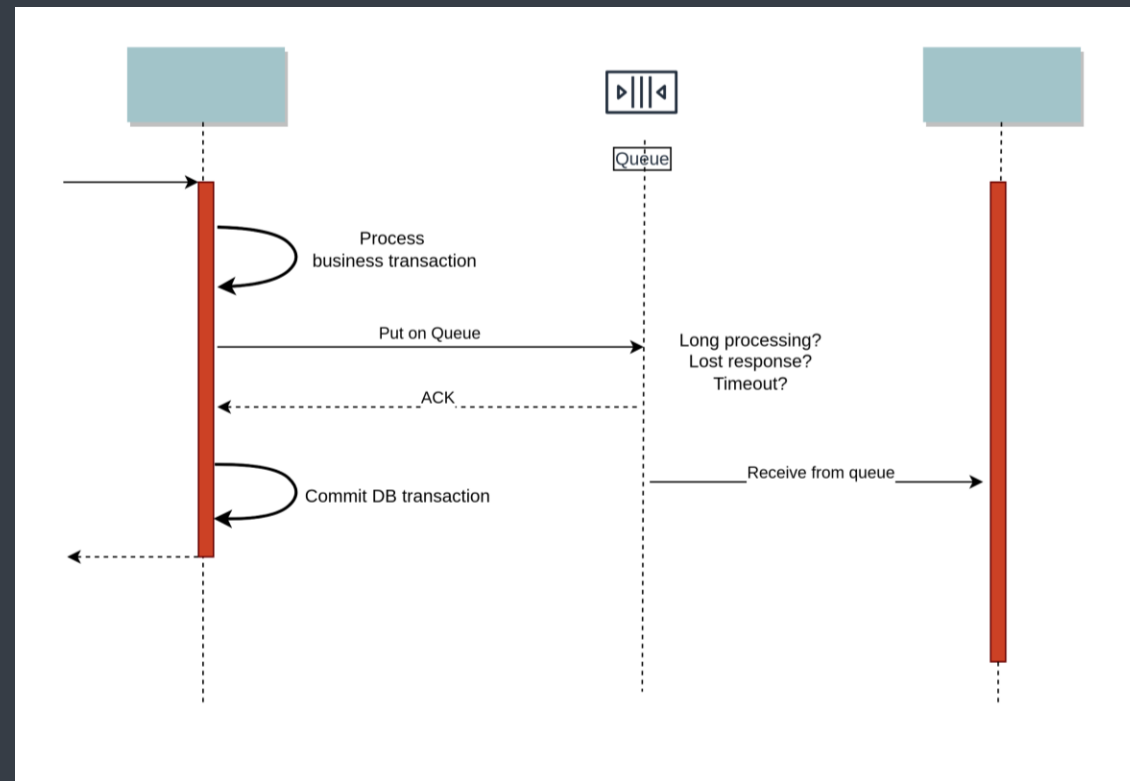


OUTBOX PATTERN

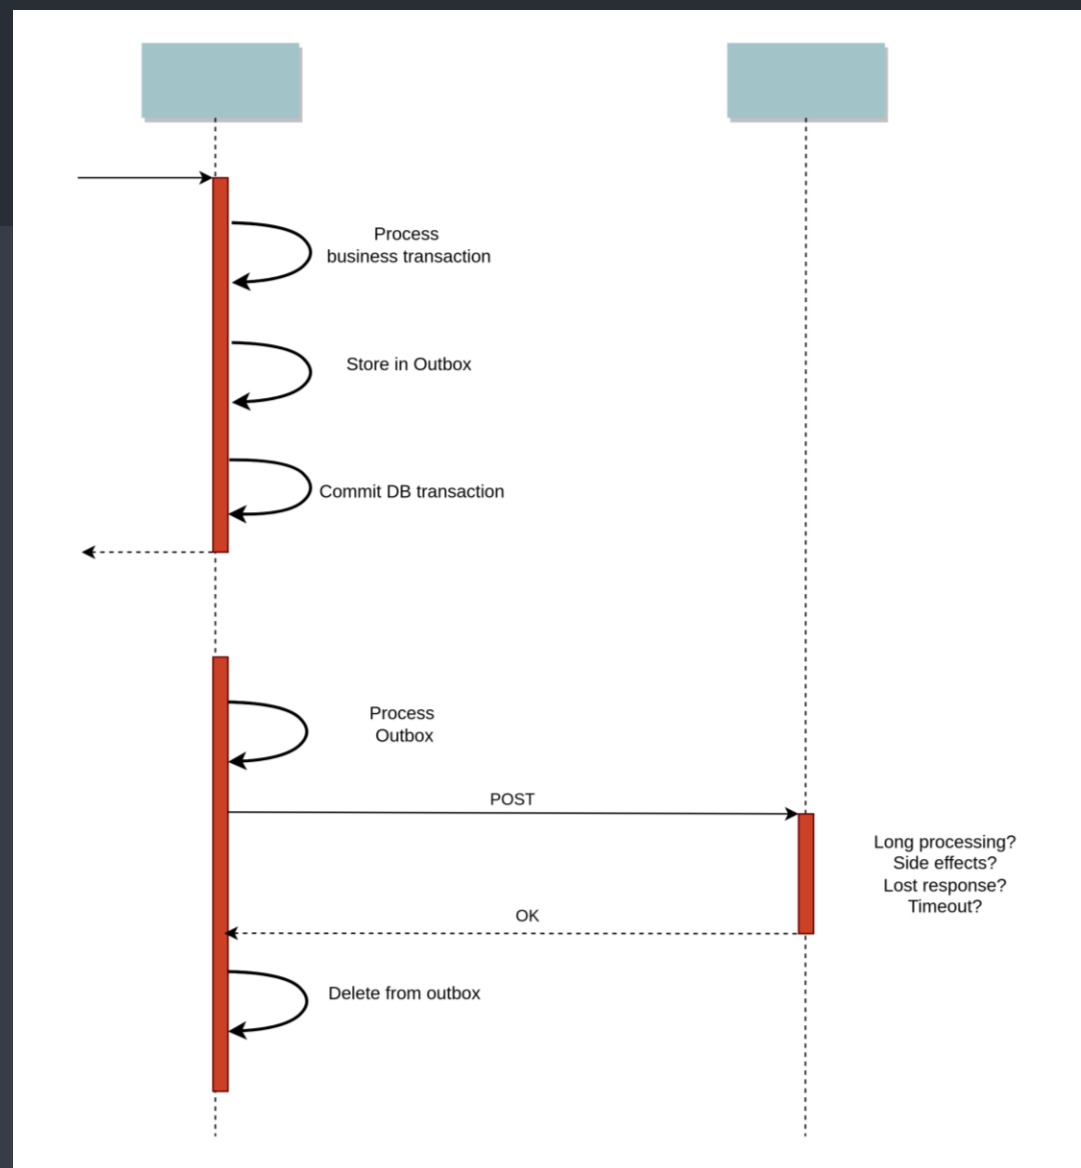
SYNC



NAIVE ASYNC



OUTBOX



OUTBOX

async

consistency

eventual

Implementation effort

Complex

Processing time

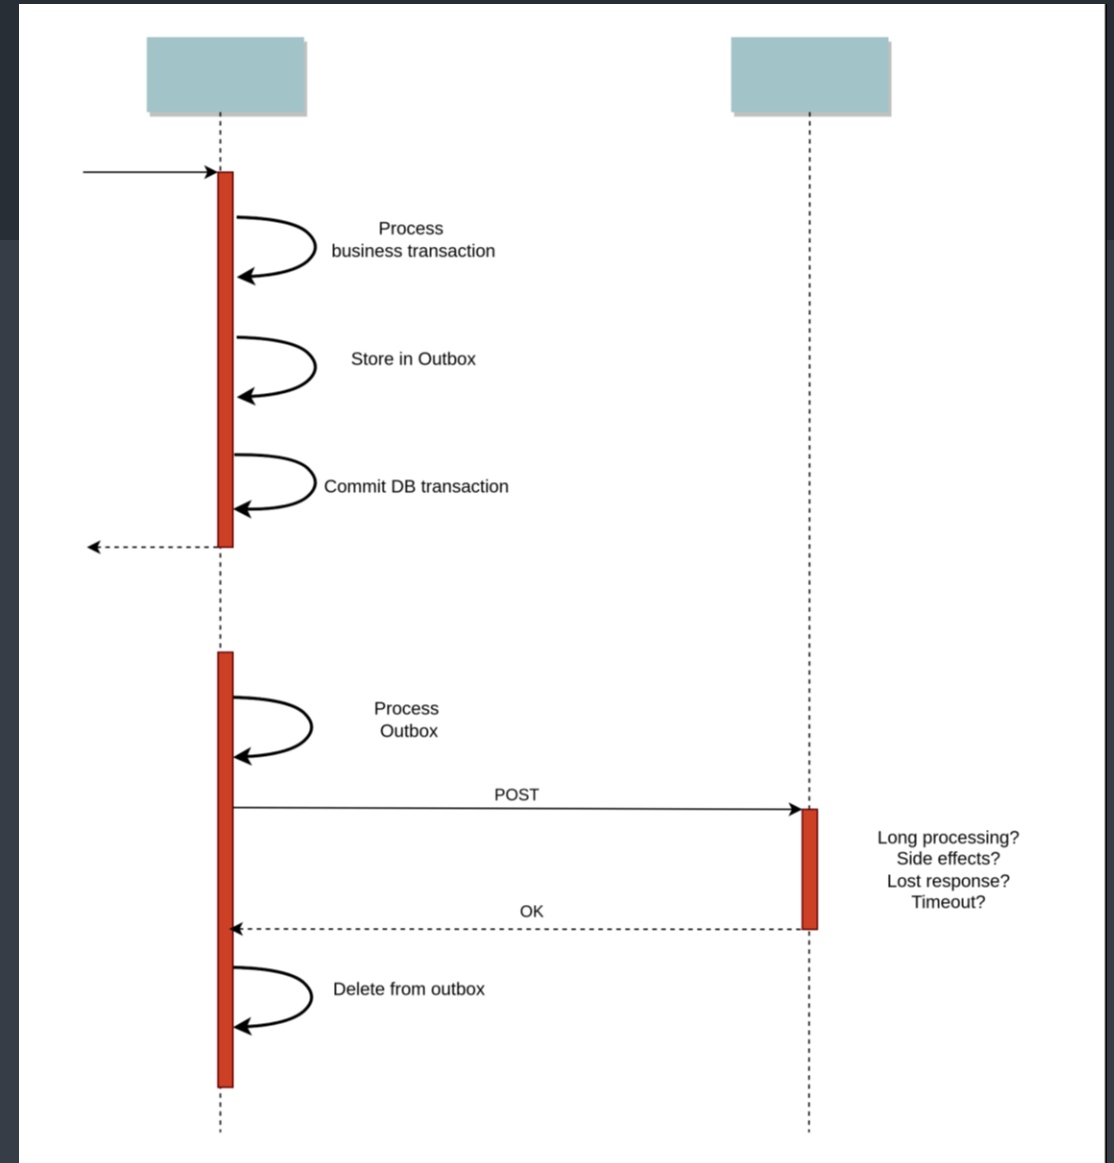
Same but split to steps

Handling errors /
rollbacks


Reconciliation needed

Retry mechanism

Simpler, depending on a
tools



OUTBOX WITH QUEUE



IS THE QUEUE NEEDED?

- HANDLES RETRIES
 - DLQ MECHANISM
- OFFLOADS DB TRANSACTIONS
 - WHEN API REQUESTS ARE SLOW
 - REDUCES THE OUTBOX TABLE SIZE
- ADDS COMPLEXITY

WHAT DOES IT SERVE


- ASYNC
- MUST BE CONSISTENT, CAN BE EVENTUALLY CONSISTENT
- REQUEST — RESPONSE
 - WITH RESPONSE IS HANDLED ASYNCHRONOUSLY
- DISTRIBUTED TRANSACTION REPLACEMENT
- AT LEAST ONCE DELIVERY

IMPLEMENTATION HOW-TO

OUTBOX EVENTS PROCESSING


SAVING TO OUTBOX

- WITHIN BUSINESS TRANSACTION
- INCLUDE ALL DETAILS NEEDED AT LATER STEPS
- ORDER
 - UNORDERED
 - GLOBAL ORDER
 - LOCAL ORDER




READING OUTBOX

- PERIODICALLY
 - SHEDLOCK, QUARTZ, CRON
 - READER ON SINGLE INSTANCE IN MULTI INSTANCE ENVIRONMENT
 - CHOOSE READING INTERVAL
- AFTER TRANSACTION COMMIT
 - REDUCING DELAY
 - BY A SEPARATE THREAD



SAVING TO QUEUE

- EACH OUTBOX EVENT PROCESSED IN A SEPARATE TRANSACTION
- MIND THE ORDERING
- OFFLOAD DB FROM HEAVY PROCESSING
- AFTER ACK:
 - REMOVE FROM OUTBOX
 - OR UPDATE ITS STATUS

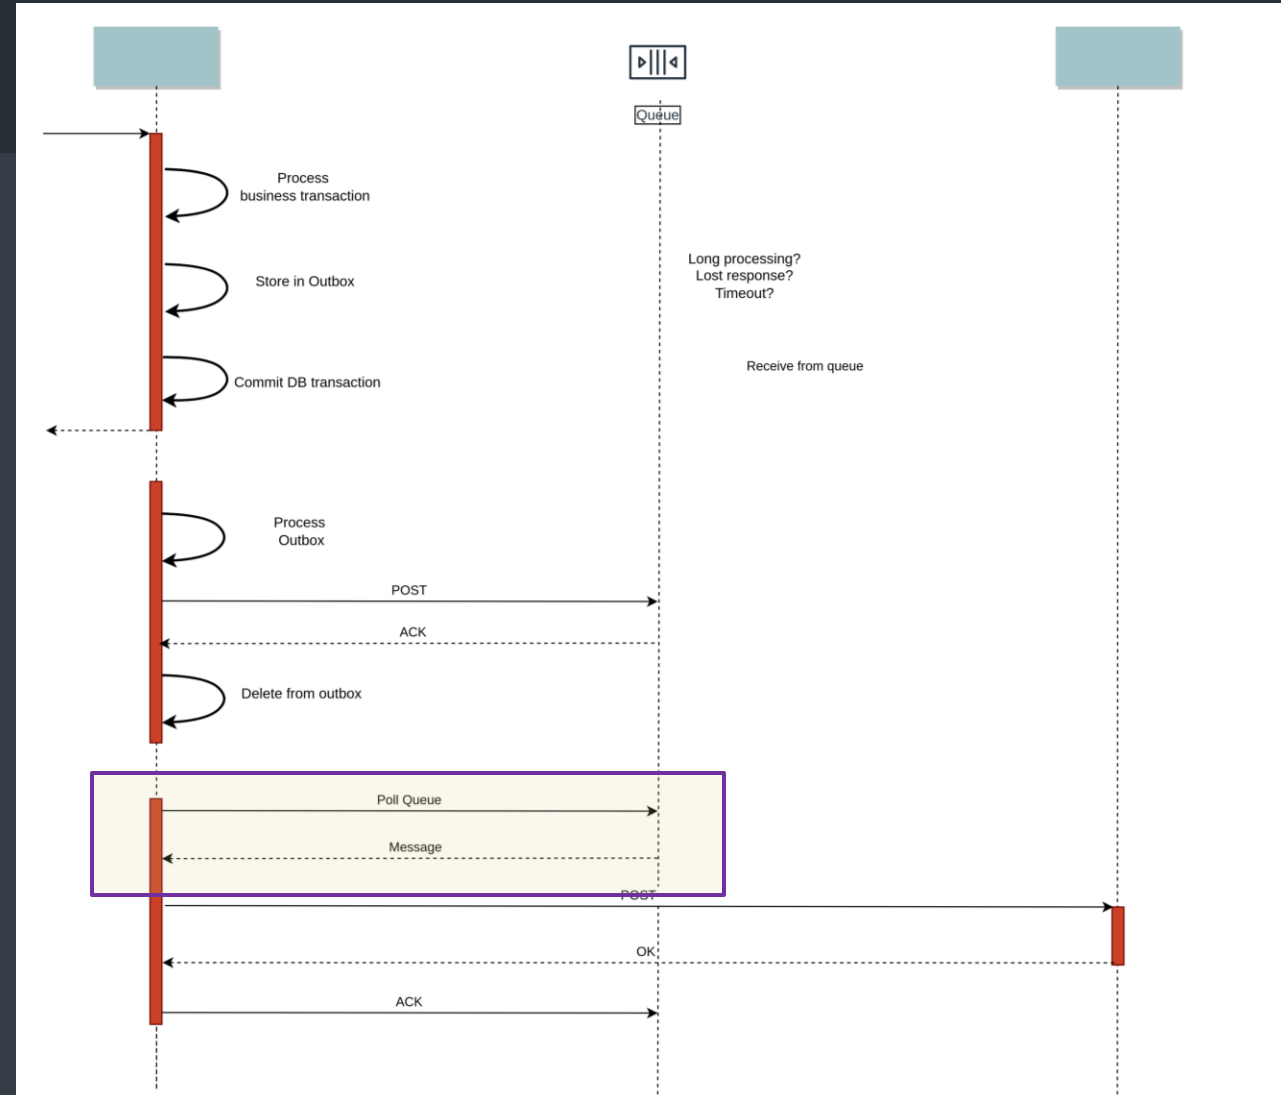


QUEUE CONFIG

- FIFO
- CONTENT BASED DEDUPLICATION
- VISIBILITY TIMEOUT: 60s
- MAX RECEIVE COUNT: 10

READING QUEUE

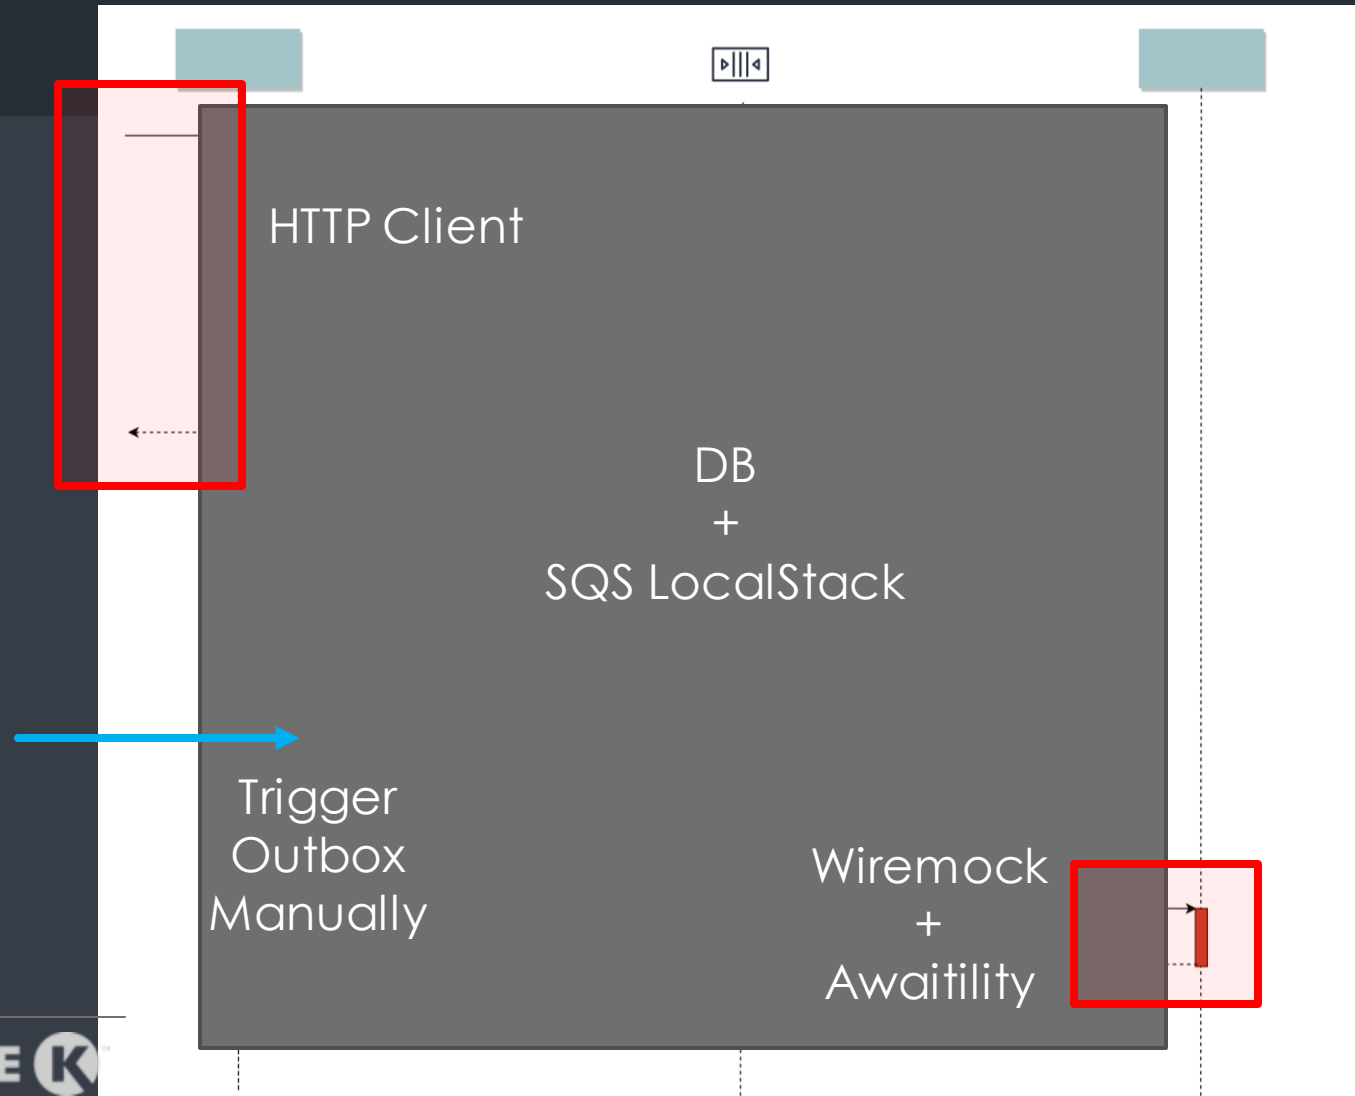
- ORDER
 - MULTIPLE THREADS MIGHT BREAK THE ORDER
- RETRY MECHANISM ON FAILED PROCESSING
 - MESSAGE COMES BACK TO QUEUE WITH VISIBILITY TIMEOUT
 - WHEN FAILED AFTER MAX RETRIES, GOES TO DLQ
- DEAD LETTER QUEUE
 - DLQ EVENTS ARE MONITORED
 - DLQ SUBSCRIBER NOTIFIES CUSTOMER CARE ON MANUAL INTERVENTION REQUIRED
- WHEN API RESPONDS WITH OUR ERROR (4xx)
 - MESSAGE LANDS IN DLQ DIRECTLY



IMPLEMENTATION HOW-TO


TESTING

INTEGRATION TESTS



UNIT TESTING

- TESTING BUSINESS AND PROCES LOGIC
- NOT VERY SPECIFIC TO OUTBOX




CROSS CUTTING CONCERNS

OUTBOX DATABASE

USE SAME DB AS YOUR PRODUCTION DB

- SQL
 - EVENT IS STORED AS A JSONB TYPE IN PSQL
- NoSQL
 - TRANSACTIONS ARE DIFFERENT BUT THE IDEA OF CONSISTENCY IS THE SAME



VERSIONING EVENTS

OUTBOX AND QUEUE EVENT IS A CONTRACT


- OLD APP VERSION MIGHT NOT BE ABLE TO PROCESS NEW VERSION
- WHEN CHANGING THE CONTRACT
 - DO BACKWARD COMPATIBLE CHANGES
 - VERSIONING IN EVENTS

```
{  
  name: "Jacek"  
}  
  
↓  
  
{  
  name: "Jacek",  
  status: "BLOCK"  
}
```

POISONING MESSAGE

BROKEN MESSAGE AT THE HEAD OF THE QUEUE

- ENSURE IT DOES NOT BLOCK THE WHOLE PROCESSING, BUT JUST ITSELF
 - IGNORE ERROR AND CONTINUE
- MIGHT HAPPEN WHEN:
 - BREAKING EVENT SYNTAX COMPATIBILITY: POSTPONE IT
 - BUG: DROP IT TO DLQ



OBSERVABILITY

- LOGGING ALL PROCES STEPS
- MONITORING THE METRICS
- TRACING ALL STEPS
 - TRACEID IS PERSISTED THROUGH THE WHOLE PROCESSING

RETRY MECHANISM

AT LEAST ONCE DELIVERY

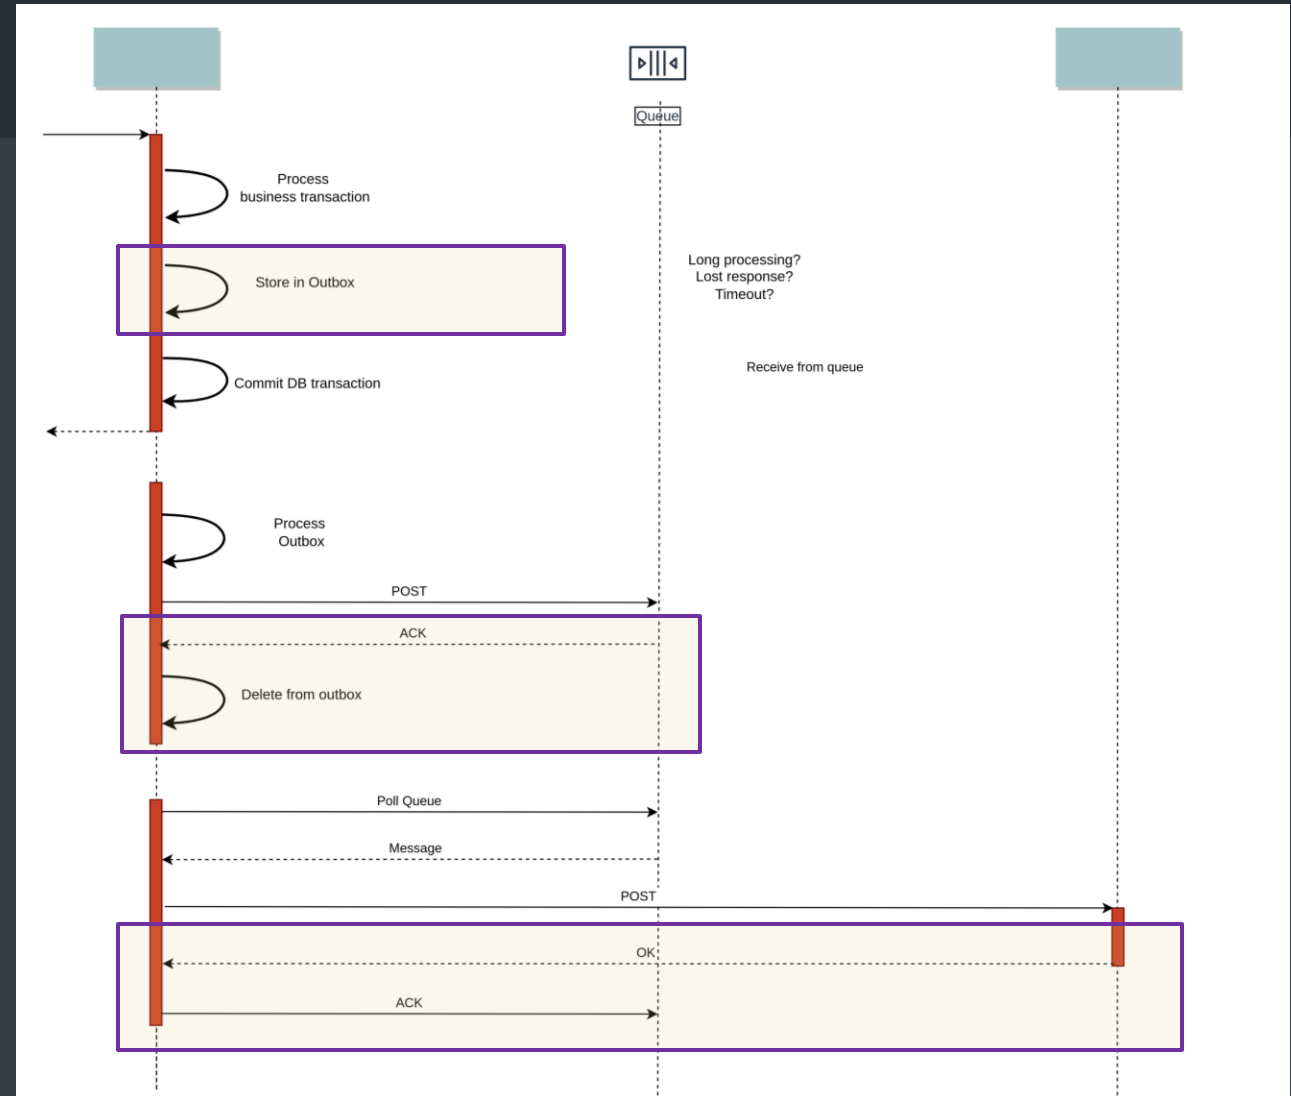
IN TWO FLAVORS:

- RETRY READING FROM OUTBOX
- RETRY WITH VISIBILITY TIMEOUT ON SQS

AT LEAST ONCE DELIVERY

DELIVERY GUARANTEED ON EACH STEP

- EVENT IS DELETED ONLY AFTER ACK IS RECEIVED
- DUPLICATION HAPPENS WHEN:
 - ACK WAS LOST
 - DELETE FAILED
 - ON SPRINT DEMO



ORDER

ORDERING NEEDS SPECIAL TREATMENT IN ASYNC COMMUNICATION

- HOW STRICT ORDER IS REQUIRED?
 - UNORDERED
 - GLOBAL ORDER
 - LOCAL ORDER
 - ORDER OF DELIVERY OR A PRESENTATION ORDER
- EVENTS NEED TO HAVE SEQUENCE OR TIMESTAMP
 - SOURCED FROM CENTRAL SOURCE OF TRUTH, E.G. DATABASE
 - READING FROM OUTBOX AND SQS HONORS THE ORDER
 - ONE THREAD AT A TIME CAN READ OUTBOX AND PROCESS EVENTS IN AN ORDERING BOUNDARIES

OUTBOX PATTERN:
WHEN SIMPLE API CALL IS
NOT ENOUGH

THANK YOU!

JACEK MILEWSKI

✉ jacek.milewski.k@gmail.com

🐦 @jacek_mil

